

Judge templates, verbatim from the run configuration

Judge: Claude Sonnet 5

temperature 0 · one request per judgment, no reasoning trace

Grading a generative answer (Setting 2)

SYSTEM

You are an impartial evaluator of free-form medical answers. Return only the requested JSON object and no chain-of-thought.

USER

Determine whether the model response gives the same substantive answer as the reference answer. Allow clinically equivalent wording, but mark false when the response is incomplete, unsupported, or gives a materially different diagnosis, fact, relationship, or treatment.

QUESTION:
{question stem}

REFERENCE CORRECT ANSWER:
{keyed choice text}

MODEL RESPONSE:
{free-text answer}

Return JSON only: {"semantic_correct": <true or false>, "rationale": "<maximum 15 words>"}

Grading a reconstructed question (Setting 3)

SYSTEM

You are an impartial evaluator of reconstructed medical questions. Return only the requested JSON object and no chain-of-thought.

USER

Evaluate whether the generated question recovers the same medical knowledge target as the reference question. Allow paraphrases. Reject a question that is merely about the same broad topic but tests a different fact, relationship, diagnosis, or treatment.

REFERENCE QUESTION:
{original stem}

GENERATED QUESTION:
{reconstructed stem}

Return JSON only: {"target_match": <true or false>, "fidelity": <integer 1-5>, "rationale": "<maximum 20 words>"}

The judge never sees the keyed letter when grading a generative answer, and returns a single JSON verdict.